

Multimodal Fashion & Context Retrieval

Glance ML Internship Assignment — Technical Submission

This document presents (1) candidate approaches and their trade-offs, (2) the chosen architecture and why it handles fashion queries better than vanilla CLIP, (3) the codebase, and (4) future-work directions for locations/weather and higher precision. Results come from two runs of the same codebase: a 648-image synthetic dataset (with explicit compositional hard negatives) used for quantitative recall, and the real FashionPedia test split (3,200 images) used for an end-to-end qualitative demo.

1. Approaches Considered (trade-offs)

Approach	What it is	Good when	Weakness
Vanilla CLIP (global embedding + cosine)	Encode image & text into one shared vector	Zero-shot; trained by both the standard cosine/similarity	Single-purpose/fails on compositionality (red shirt+blue pants ≈ blue shirt)
Attribute classifiers + keyword index	Detect colour/garment/scene with late-interaction	Interpretable exact tags, attribute-level recall	Needs labelled training data or a fine-tuned detector; brittle to novel combinations
Late-interaction (CLIP + attribute re-rank)	CLIP for recall; a disentangled (co-occurrence) model for re-ranking	Keeps CLIP's zero-shot ability while adding specificity	Quality bounded by the attribute model
Specialised compositional models (TAI/CoCa, BLIP-2)	Train/fine-tune a model that explicitly handles composition	Hybrid attributes	Heavy, needs fashion training data/GPU; overkill for a retrieval task
Text-to-image generation as retrieval (reverse)	Generate the query image, then match	Makes novel descriptions	Slow, unstable, not exact.

Selection rationale: the brief asks to spend effort on **ML logic**, not on re-implementing a vector DB, and explicitly wants something **better than vanilla CLIP with a fashion focus**. The late-interaction hybrid delivers exactly that with the least engineering risk and stays zero-shot.

2. Chosen Approach — Attribute-Aware Compositional Retrieval

Pipeline: **query** → **parser** → **indexer (FAISS + attribute tables)** → **retriever (recall then compositional re-rank)**.

2.1 Query understanding

A parser decomposes free text into a structured `ParsedQuery{scene, style, bindings[(color, garment)...], free_text}`. Crucially, colour and garment are kept as **paired bindings**, not a flat word list. A rule-based parser maps the free text onto the taxonomy with no network calls.

2.2 Indexer (Part A)

- **Feature extraction:** each image produces a global CLIP embedding *and* a disentangled attribute table — a logit for every (colour, garment) pair, plus scene and style logits (zero-shot prompts like “a photo of a red shirt”).
- **Vector storage:** FAISS IndexFlatIP (exact cosine) holds the global embeddings for fast recall; the attribute tables live in a side metadata store. No filename/keyword matching.
- **Scalability:** FAISS IVF/PQ (or Pinecone/Milvus) swaps in for 1M+ images with no change to the retriever; only the candidate set grows.

2.3 Retriever (Part B) — the compositionality fix

FAISS returns the top-N candidates by global similarity (cheap recall). The final ranking is decided by a **compositional score**:

$$\text{final} = w_g \cdot \text{global} + w_a \cdot \text{attr} + w_s \cdot \text{scene} + w_t \cdot \text{style}$$

The attribute term enforces the binding: for each query binding (colour c, garment g) we score the image's *specific* (c, g) cell and normalise it against the strongest colour predicted for that garment,

$$s_b = \text{attr}[(c,g)] / (\max_{c'} \text{attr}[(c',g)] + \epsilon)$$

so the image must have the **shirt be red**, not merely contain red somewhere and a shirt somewhere. Swapped bindings therefore score ≈ 0 and are pushed down. This is what vanilla CLIP cannot do.

2.4 Zero-shot & fashion focus

No fine-tuning is required: the attribute logits and the parser both operate over an open taxonomy, so unseen colour/garment descriptions are handled at inference. The taxonomy is fashion-specific (blazer, hoodie, raincoat, tie ...) which lifts fine-grained fashion accuracy over a generic CLIP prompt set.

3. Experimental Results

Dataset: 648 synthetic images spanning 4 scenes $\times$ multiple garment types $\times$ 16 colours, including balanced **compositional hard negatives** (e.g. both red-shirt/blue-pants and blue-shirt/red-pants).

3.1 Assessment-query recall (vanilla vs compositional)

Query	#rel	vR@1	vR@3	vR@5	vR@10	cR@1	cR@3	cR@5	cR@10
A person in a bright yellow rainco	10	0.10	0.20	0.30	0.50	0.10	0.30	0.50	1.00
Professional business attire insid	47	0.02	0.06	0.11	0.21	0.02	0.06	0.11	0.21
Someone wearing a blue shirt sitti	17	0.06	0.18	0.24	0.35	0.06	0.18	0.29	0.59
Casual weekend outfit for a city w	41	0.02	0.07	0.12	0.24	0.02	0.07	0.12	0.24
A red tie and a white shirt in a f	13	0.08	0.08	0.15	0.38	0.08	0.23	0.38	0.69

Mean R@10: vanilla 0.34 $\rightarrow$ compositional 0.55. Queries 2 & 4 are style/scene-only (no bindings) and sit near their recall ceiling because many images satisfy them; the lift is concentrated where fine-grained / compositional constraints apply (queries 1, 3, 5).

3.2 Compositionality binding probe (lower rank = better)

For each colour $\leftrightarrow$ garment swap pair we ask for the correctly-bound outfit and compare the mean rank of the **correct** image vs the **swapped** hard-negative image.

Query	Mode	rank(correct)	rank(swapped)
a red shirt and a blue pants	vanilla	25.9	31.1
a red shirt and a blue pants	compositional	11.5	86.3
a blue shirt and a red pants	vanilla	31.1	25.9
a blue shirt and a red pants	compositional	12.5	67.3
a white shirt and a black pants	vanilla	28.2	27.9
a white shirt and a black pants	compositional	13.0	76.2
a black shirt and a white pants	vanilla	27.9	28.2
a black shirt and a white pants	compositional	10.0	52.6
a red tie and a white shirt	vanilla	28.9	29.1
a red tie and a white shirt	compositional	13.5	82.8
a white tie and a red shirt	vanilla	29.1	28.9
a white tie and a red shirt	compositional	15.5	74.6

Reading: vanilla CLIP ranks correct and swapped images almost identically (~ 26 – 31), i.e. it cannot tell them apart. The compositional retriever pulls the correctly-bound image to rank ~ 10 – 15 while pushing the swapped hard negative to rank ~ 52 – 86 . This is a direct, quantitative demonstration of the compositionality fix.

3.3 Real dataset — FashionPedia test split (3,200 imgs)

The same pipeline runs unchanged on the real FashionPedia test split (val_test2020.zip, 3,200 images, no labels). A CLIP (ViT-B/32) index was built; the five assessment queries return the retrieved images below. On real photos the

compositional re-rank moves little versus vanilla CLIP, because the attribute signal comes from zero-shot CLIP, which is known to be weak on fine-grained colour/garment classification. That is precisely the limitation the brief calls out: the fix is to replace the zero-shot attribute table with a fine-tuned fashion detector or BLIP-generated captions, after which the same re-rank recovers the gain shown in 3.1/3.2. The architecture does not change.

[examples not rendered: ... must have at least a row and column]

4. Codebase

GitHub: <https://github.com//glance-fashion-retrieval> (push the repository root; modular layout:)

src/ — taxonomy, query_parser, feature_extractor (Synthetic + CLIP backends), store (FAISS), indexer, retriever, data_generator
scripts/ — build_index.py, search.py
eval/ — run_evaluation.py, results.json
data/, index/ — generated at run time

Reproduce:

```
pip install -r requirements.txt
python scripts/build_index.py --data data --index index --n 600
python eval/run_evaluation.py
python scripts/search.py "A red tie and a white shirt in a formal setting"

# real FashionPedia split (needs torch + open_clip_torch)
python scripts/build_index.py --images-dir data/fashion2020/test --index index_real --backend clip
python eval/run_real.py
python scripts/search.py "A red tie and a white shirt in a formal setting" --index index_real --clip
```

5. Future Work

5.1 Extending to locations (cities, places) and weather

- **Geo / city:** add a *location* axis to the taxonomy and a geo-embedding (e.g. a place recognition backbone or GPS metadata). Store location as an extra attribute table; the parser extracts city/landmark mentions (LLM or gazetteer). Query becomes {scene, location, style, bindings}.
- **Weather:** add a *weather* axis (rain/snow/sun) detected from scene pixels or an external signal; bind it to outfit appropriateness (raincoat↔rain). The same attribute re-rank naturally absorbs it.
- **Multi-signal fusion:** learn the weight vector w from click/log data instead of hand-tuning, so location/weather importance adapts per query.

5.2 Improving precision

- **Better attribute tagger:** replace zero-shot prompts with a fine-tuned fashion attribute model (e.g. on FashionPedia) for sharper (colour, garment) logits.
- **Cross-modal late interaction (TAI-style):** keep token-level image region ↔ text token maxima so fined-grained regions bind to words.
- **Hard-negative mining + re-ranking loss:** train a small adapter on triplets (query, correct, swapped) to push the binding score margin up.
- **Calibration & thresholds:** per-axis sigmoid calibration plus a learned combination (linear/GBM) over the four scores using labelled relevance.
- **Quantisation + IVF/PQ:** for 1M+ images use FAISS IVF-PQ so recall stays sub-ms while the re-rank runs only on top-N.

Summary: the submission is a modular, zero-shot, attribute-aware retriever that demonstrably fixes CLIP's compositionality gap through explicit (colour→garment) binding, scored on a disentangled attribute table and re-ranked over FAISS recall.